Firstake Prompt Rules

Oct 26, 2026 · @zara

How Firstake writes a prompt

Every request runs through the same seven steps, whatever the media type. The user sees only the result of step 7; steps 1–6 happen behind the scenes.

1. **Classify the job.** Image, consistent character, video, music, voice, website, slides, or a light category (writing, code, automation). One request can be several jobs, such as a promo video that also needs a voiceover.
2. **Check the slots.** Each media type has required slots (listed in its section below). Fill as many as possible from the user's words.
3. **Ask only for what's missing.** If every required slot is filled, skip straight to step 4. Otherwise ask at most two short questions, the ones that change the result the most, in plain words ("how wide should it be?" not "aspect ratio?"). Optional slots get sensible defaults instead of questions.
4. **Pick the tool.** Use the catalog's ranking: fit first, then cost per finished output. Show the top pick, the cheapest good pick and the free option, with prices.
5. **Fill the template** for that tool, in the syntax it understands (sentences for GPT Image, tags for Suno, parameters for Midjourney, shot lists for Veo).
6. **Add the avoid-list.** Name the generic defaults that tool falls into for this kind of job, and steer away with specific choices, not just bans.
7. **Hand it over.** The ready-to-paste prompt, one line on why this tool, one mistake to avoid, and for beginners, where to go and where to paste it.

Matching the user's level: if the user writes in plain language, answer in plain language with click-by-click steps. If they use tool jargon ("9:16", "--sref", "dolly-in"), answer briefly and technically.

Images

The biggest single upgrade to any image prompt is replacing vague quality words ("realistic", "beautiful", "high quality") with the concrete things that create that look: a camera, an era, a light source, real-world details.

Slots

Slot	Required?	Ask in plain words	Default if not given
Subject	Yes	"What should be in the picture?"	—
Purpose / where it's used	Yes	"Where will you post or print it?"	Sets size and style
Look (photo, illustration, 3D, painting…)	Yes	"Should it look like a real photo or a drawing?"	Photo
Exact text in the image	If any	"Any words that must appear, spelled exactly?"	None
Shape (aspect ratio)	No	"Tall, wide or square?"	From purpose: 9:16 story/reel, 4:5 Instagram post, 16:9 slide/YouTube, 1:1 profile
Mood, light, colors	No	—	Matched to purpose
Reference image	No	"Do you have a photo to start from?"	None

Universal template (filled in this order):

```
[Look + camera/medium], [subject doing something specific], [setting with 3–5 real details], [light source and time of day], [color/mood], [composition and shape]. Text: "[exact words]". [What it must not look like, phrased as what it should look like instead].
```

Rules that apply to every image tool

- **Realism comes from cameras and eras, not adjectives.** "Amateur iPhone photo", "shot on 35mm film", "direct on-camera flash", "2003 digital camera snapshot" each produce a believable look that "photorealistic" does not (awesome-gpt-image).
- **Ordinary details beat polish.** Naming small real-world objects (stickers on a fridge, a shopping cart, water drops on the ground) and saying people should look like ordinary passers-by keeps images from looking staged.

- **Every word of text goes in quotes, line by line.** For posters and ads, list each headline, price and fine-print line separately, then say which matters most.
- **Short prompts for simple ideas, structured prompts for complex ones.** A one-line idea ("90s point-and-shoot photo of…") is fine; a product ad with ten text lines needs labeled parts.
- **Say the anti-goal positively.** "Premium and calm, not cluttered" works better than a long list of "no" words.
- **No real people or brand logos by default.** Community prompt libraries are full of celebrity and brand images; Firstake should not write those, and should warn that such images can't be used commercially.

Tool-specific rules

Tool	Write it like this	Key controls	Common mistake
GPT Image 2 (ChatGPT)	Full sentences, like briefing a designer; follows long, detailed instructions well	Exact text in quotes; upload a photo and say what to change	Asking for a style without naming the medium
Nano Banana Pro / 2 (Gemini)	Conversational; describe the scene as a story, not a keyword list	Edit in steps: "keep everything the same, only change X"; combine several uploaded images	Re-describing the whole image when you only want one change
Midjourney V8	Short descriptive phrases, most important first	`--ar` shape, `--stylize` artistic strength, `--raw` less stylized, `--sref` style reference, `--oref` + `--ow` character/object reference, `--no` to exclude, `--seed` repeatability	Expecting the same face from the same text (see Consistent characters)
FLUX 2	Natural language with camera and lens terms ("85mm, f/1.8, shallow depth of field")	Reference images for consistency	Long lists of negative words

Tool	Write it like this	Key controls	Common mistake
Ideogram 3.0	Name the layout first ("poster with headline at top"), then the text in quotes	Style presets; magic-prompt toggle rewrites your prompt, so turn it off when wording is exact	Misspellings in the prompt get copied into the image
Adobe Firefly	Plain description plus style presets	Content type (photo/art), style and composition reference images	Expecting edgy styles; it is tuned to be commercially safe
Leonardo	Pick the model and preset first, then describe	Character reference, image guidance strength	Using a mismatched model for the style

Avoid-list for images (Firstake steers away by naming specifics): plastic-smooth skin, everything perfectly centered, generic blue-orange "cinematic" grading, the same glossy stock-photo look, extra fingers on close-up hands, garbled background text.

Consistent characters

The same text prompt never gives the same character twice, because every image starts from a new random seed. Consistency comes from a saved reference image, not from better wording. Firstake switches to this workflow whenever the user says "same character", "my mascot", "comic", "series", "keep her looking the same", or complains that a character keeps changing.

Triggers and the one question to ask: "Do you already have a picture of the character you like?" If yes, start at step 3. If no, start at step 1.

The anchor workflow (all tools)

1. **Design once, on purpose.** Generate the character on a plain background, full body, neutral pose, with a written "character card": hair, eyes, build, outfit, colors, one or two distinctive features.
2. **Make a reference sheet.** Ask for front, side and back views plus a few expressions in one image, with "face, hairstyle and clothing must stay identical in every panel" (awesome-gpt-image). Save it as the anchor.

3. **Attach the anchor every time** and paste the same character card word for word. Only the scene and action change.
4. **Change one thing at a time.** New pose or new place, not both plus a new outfit.
5. **Fix drift by strengthening the reference, not rewriting the prompt.**

Per tool

Tool	How to lock the character	Strength setting	Watch out for
Midjourney	`--oref [anchor image link]`; add `--sref [same link]` to lock the art style too	`--ow` 1–1000, default 100; raise toward ~400 if the face slips; above ~400 it copies the photo	Costs 2x GPU time; the old `--cref` is replaced; Midjourney video has no `--oref`
Nano Banana	Upload the anchor and say "keep this exact character, change only the scene to..."	Keep the instruction short; re-upload the anchor in long chats	Long edit chains slowly drift; restart from the anchor
GPT Image 2	Upload the anchor; describe the new scene; repeat the character card	—	Same drift risk over long chats
Veo 3.1 (video)	"Ingredients to video": one reference image per character and one for the setting, referenced by role ("using the provided images for the detective...")	—	Generate the ingredient images first in Nano Banana (Google Cloud)
Leonardo / OpenArt	Built-in character reference tools that save a character for reuse	Reference strength slider	Credits per generation are higher

Tool	How to lock the character	Strength setting	Watch out for
Many images, near-perfect match	Train a small custom model (LoRA) on 15–30 images of the character	—	Needs setup; best for comics, mascots, brands

Example prompt Firstake hands back (Midjourney)

```
[character card, pasted identically], sitting at a café table reading, morning window light --oref [anchor link] --ow 150 --sref [anchor link] --ar 4:5
```

This is a strong demo case: a user says "I made a character in Midjourney but it keeps changing", and Firstake explains why and fixes it in one answer.

Video

A video prompt is a shot description, not a story. Each generation is one short scene (4–10 seconds), so Firstake breaks longer ideas into a list of shots and writes one prompt per shot.

Slots

Slot	Required?	Ask in plain words	Default
What happens	Yes	"What should happen in the clip?"	—
Where it's posted	Yes	"Is this for TikTok/Reels, YouTube, or a presentation?"	Sets 9:16 or 16:9
Length	Yes	"About how many seconds?"	8 s; longer = several shots
Starting image	No	"Do you have a photo to bring to life?"	Text only
Sound or dialogue	No	"Should it have talking, sound effects, or be silent?"	Ambient sound if the tool supports it
Camera, light, mood	No	—	Chosen to fit the purpose

Universal formula (Google's official Veo order, which works across most tools):

```
[Camera: shot type + movement] + [Subject] + [Action, as physical motion] +
[Setting] + [Style, light and mood] + [Audio]
```

Example: *Medium shot, slow push-in. A baker in a flour-dusted apron pulls a tray of croissants from the oven, steam rising. Small warm bakery at 6 a.m., golden window light, cozy and calm. SFX: oven door creak, soft clink of the tray. Ambient noise: quiet morning street.*

Rules that apply to every video tool

- **Image-to-video is the beginner's best path.** A still image locks the look, subject and framing, so the video model only has to add motion and fails far less often. Firstake should usually route: make the image first (section Images), then animate it.
- **When a starting image is used, describe only motion.** Re-describing what's already in the image can reduce movement or cause odd results (Runway Help). Write "The man extends his arm and nods", not "The tall man with black hair in a gray suit extends...".
- **Turn feelings into actions.** "The woman smiles and waves" works; "a joyful greeting full of warmth" does not.
- **Say what you want, not what you don't.** "Locked camera, the camera stays still" instead of "no camera movement"; "a desolate landscape with no roads" instead of "no man-made structures" (Google Cloud). Some tools flip negative words into the

opposite result.

- **One scene per clip.** Don't combine a transformation, a season change and a spinning camera in one generation.
- **No conversational requests.** "Can you please add my dog" wastes words; write "A dog runs into the frame from the left".
- **Test cheap, finish expensive.** Draft at 720p or on a cheap model (Kling daily credits, Veo 3.1 Lite), then re-run the winner at full quality.

Camera vocabulary Firstake translates into (users can say it plainly; Firstake writes the term)

User says	Firstake writes
"zoom in slowly"	slow push-in / dolly in
"follow them"	tracking shot following the subject
"from above"	high-angle / aerial / overhead shot
"spin around them"	180-degree arc shot
"keep it still"	locked camera, static shot
"blurry background"	shallow depth of field
"like filmed on a phone"	handheld, slight shake, vertical

Tool-specific rules

Tool	Write it like this	Special features	Common mistake
Veo 3.1	The five-part formula above	Dialogue in quotes (`A woman says, "We have to leave now."`); `SFX:` and `Ambient noise:` lines; timestamped multi-shot `[00:00-00:02]` ; first + last frame; ingredients (reference images); 4, 6 or 8 s; 16:9 or 9:16	Writing a whole story in one 8-second clip

Tool	Write it like this	Special features	Common mistake
Runway Gen-4.5	Start simple with the main motion, then add one element at a time: subject motion, camera motion, scene motion, style	Refer to subjects as "the subject", "she", or "the man on the left"	Negative phrasing; over-describing the input image
Kling 3.0	Subject + action + camera + pace words (slowly, smoothly, suddenly)	Strong physics: liquids, cloth, collisions; separate negative-prompt box	Asking for long multi-scene clips (max ~10 s)
Hailuo	Sequenced steps ("the camera descends, then circles right")	Strong facial expressions and emotion changes	Vague camera direction
Seedance 2.0	Natural language with camera terms	Cheap per second; audio and video together	Expecting top photorealism
Midjourney Video	Animate an existing Midjourney image; choose low or high motion	Keeps Midjourney's art style	No character reference in video

Avoid-list for video: warped faces and hands in close-ups, objects melting into each other, unreadable on-screen text (add text later in an editor), sudden style changes mid-clip, too many moving things at once.

Music and voice

Music and voice tools split instructions across separate boxes, and most bad results come from putting the right words in the wrong box. Firstake's output for these tools is not one prompt but a filled-in form: one block per field.

Slots

Slot	Required?	Ask in plain words	Default
Use	Yes	"Is it for fun, or for a business, ad or monetized channel?"	Personal → free tiers allowed
Genre and mood	Yes	"What kind of music, and how should it feel?"	—
Vocals or instrumental	Yes	"Singing, or just music?"	—
Lyrics	If vocals	"Do you have lyrics, or should I write them?"	Firstake drafts them with a cheap LLM
Length, tempo	No	—	Tool default
Voiceover script, voice type	For voice jobs	"What should it say, and what kind of voice?"	—

Licensing check first: free plans on Suno and Udio are personal use only. If the use is commercial, Firstake routes to a paid Suno plan or ElevenLabs and says so up front.

Suno: what goes in which box (Jack Righteous guide)

Box	Put here	Don't put here
Styles	One foundation genre + a few accents, key instruments, vocal type, mood, era, optional tempo ("warm 70s soul, female vocal, Rhodes piano, slow groove, nostalgic")	Lyrics, long sentences, artist names
Lyrics	The words, with section labels: `[Intro]`, `[Verse]`, `[Pre-Chorus]`, `[Chorus]`, `[Bridge]`, `[Outro]`; short local cues like `[Bridge – solo cello]`	Prose directions, which can get sung out loud
Exclude	Unwanted sounds ("trap hi-hats", "autotune")	Positive instructions
Settings	Length, style influence, vocal gender	—

Suno rules: use Custom mode, not Simple, so problems are easy to trace. Keep chorus lines short and repeatable. When a result misses, change one thing at a time (style words, then lyrics, then settings). Describe the sound of an artist ("breathy indie-pop vocal, lo-fi guitar") instead of naming the artist. Tempo and key are suggestions, not guarantees.

ElevenLabs voice: how to direct delivery (ElevenLabs docs)

- **Audio tags in square brackets** right before the line they change: `[whispers]`, `[sighs]`, `[laughs]`, `[excited]`; sound tags like `[applause]`.
- **Describe voice quality clearly** (`[low, gravelly voice]`) so the tag isn't read as a sound effect.
- **Punctuation shapes the read:** an ellipsis (…) adds a pause, CAPITALS add emphasis.
- **Pick a voice that already fits.** A voice trained on calm narration will whisper or shout less reliably than one that naturally does.
- **Dialogue:** give each speaker a different voice and label lines ("Speaker 1:").

Example Firstake hands back for a bakery ad voiceover:

```
[warm, friendly] Fresh out of the oven… every single morning. [excited] Come taste the BEST croissants in town. [softly] We'll save you a warm one.
```

Websites and UI

AI-built sites look alike because the model fills every undecided choice with the average of the web: a centered headline, a small label above it, three icon cards, a soft gradient. Banning those defaults is not enough; the taste-skill authors found that a page stripped of them and given nothing in return looks thin and loses to a busier one (taste-skill). The fix is a specific idea taken from the subject, plus real content. That is why the website interview asks about the business, not about design.

Slots (the subject inventory)

Slot	Required?	Ask in plain words
What it is and the one action visitors should take	Yes	"What is it, and what should a visitor do: book, buy, call, sign up?"
Site or app	Yes	"Is it a page people read, or a tool people log into?" (routes to Framer/Wix vs. Lovable/Bolt)

Slot	Required?	Ask in plain words
Real content	Yes	"Paste your menu, prices, hours, services" (placeholder text produces generic layouts)
What makes it unlike its competitors	Yes	"What's one thing only you have: the name's meaning, your street, a ritual, a rule you follow?"
Existing brand	No	"Do you have a logo, colors or fonts?"
Real photos	No	"Do you have your own photos?" (never plan around pictures that don't exist)
A site you like	No	"Name a site whose look you like"

Rules Firstake builds into every website prompt

- **Find the idea with the competitor test.** Write "On this site, [something only this business has] becomes [the main visual device]". If the bakery two streets away fits the same sentence, it isn't an idea yet (taste-skill).
- **Write the words before the design.** One claim per section; a section with no claim gets cut. This avoids the stock hero → logos → three features → testimonial → pricing sequence.
- **Headlines fit in two lines** (about 32 characters), with one plain sentence under them saying what the thing is, and one clearly named button.
- **Full content, never invented proof.** Every section delivers its real information (the whole menu with prices). No made-up reviews, customer logos, press quotes or statistics.
- **One interactive thing built from the subject** (pick a pastry and see today's batch time; enter a date and see open tables), not decoration.
- **Name colors, fonts and a reference**, or a DESIGN.md file. "Make it modern" reliably produces the purple-gradient default.
- **Build one section at a time**, and give tools like Lovable a project knowledge file with the vision, users and design rules (Lovable docs).

Avoid-list (distilled from Impeccable's craft rules, impeccable): grids of same-size icon cards as the page structure; the big-number "hero metric" block; small labels above every heading; section numbers 01/02/03 that carry no meaning; gradient text; glass/blur used

as decoration; thick colored stripes on the side of cards; emoji standing in for icons; a system font as the headline face; light or dark mode picked by habit rather than by where the site is used; animations that fade in every section the same way. Plus the checks it enforces: readable contrast, body lines 65–75 characters wide, and hover, loading, error and empty states that actually exist.

Recommended stacks by user level

User	Stack	Firstake's output
Beginner, marketing site	Framer AI or Wix	Answers to the site builder's questions + section-by-section copy
Beginner, working app	Lovable	Knowledge file + first prompt + next 3 prompts, one feature each
Comfortable with code tools	Claude Code + taste-skill or Impeccable + a DESIGN.md + Playwright CLI	Install steps in plain words, the DESIGN.md, the build prompt, and a "check your own work in the browser" step

About DESIGN.md files: the awesome-design-md collection has ready files for about 75 well-known sites, each with sections for colors, typography, layout, elevation, and do's and don'ts. They're good for learning structure, but copying a famous brand's exact look onto a business site can confuse customers and invite trademark trouble. Firstake should use them as a starting point and swap in the user's own colors and fonts, or write the user a fresh DESIGN.md from their interview answers.

Slides and graphics

Slide tools build better decks from an outline than from a topic. Firstake's job is to turn "make a presentation about X" into a slide-by-slide outline before it reaches Gamma or Canva.

Slots

Slot	Required?	Ask in plain words	Default
Topic and audience	Yes	"Who will watch it?"	—

Slot	Required?	Ask in plain words	Default
Length	Yes	"How many minutes, or how many slides?"	1 slide per minute
The one takeaway	Yes	"What should people remember afterwards?"	—
Source material	No	"Do you have notes or a document to build from?"	Firstake drafts the outline
Brand colors, school or company template	No	—	Clean neutral theme
Speaking or sending	No	"Will you present it live, or send it to read?"	Live → fewer words

Rules

- **Outline first.** One line per slide: a headline that states the point ("Most people pick the wrong AI tool for the job"), not a topic label ("Problem").
- **Paste the outline, not a paragraph**, and tell the tool to keep your text rather than rewrite it when the wording matters.
- **Live talks get few words.** One idea per slide, a picture or a number, speaker notes for the rest.
- **Set the slide count yourself.** Tools pad to their default length.
- **Graphics with text** (flyers, social posts): route to Canva for layout, or Ideogram/GPT Image when the image itself must contain the words. Always give the exact wording in quotes.
- **Credits:** Gamma's free credits never refill, so Firstake sends a user's first deck there and later decks to Canva Free.

Example Firstake hands back (2-minute pitch, 5 slides)

```
1. Too many AI tools, and people pick the wrong one
2. Firstake: describe what you want, get the right tool and the right prompt
3. Demo: "promo video for my bakery" → tool, price, ready prompt
4. Neutral and cost-aware: best pick, cheapest pick, free pick
5. We need: one front-end developer and one designer
```

That last example doubles as a starting outline for the Pitch Day deck, which allows exactly two minutes and asks teams to say whether they need more people.

Sources

GitHub projects downloaded and read:

- awesome-gpt-image: GPT Image 2 and Nano Banana prompt library, side-by-side comparisons
- awesome-ai-video-prompts: beginner video guide; most tool-specific pages in this repo are still empty placeholders, so its claims were checked against official guides
- impeccable: craft rules and design anti-patterns
- taste-skill: subject-first design process and anti-default rules
- awesome-design-md: DESIGN.md format and examples
- awesome-generative-ai and awesome-ai-tools: tool directories used to find catalog gaps

Official and vendor guides:

- Google Cloud: Ultimate prompting guide for Veo 3.1
- Runway Help: Gen-4 prompting guide
- ElevenLabs: Prompting best practices
- Lovable docs: Best practices
- Jack Righteous: Where to put your Suno prompt
- Rangy: AI character consistency and UD: Midjourney consistency